



deeplearning.ai

# Face recognition

---

What is face  
recognition?

# Face recognition

A Video Here

# Face verification vs. face recognition

## → Verification

- Input image, name/ID
- Output whether the input image is that of the claimed person

1:1

99.0%

99.9

## → Recognition

- Has a database of K persons
- Get an input image
- Output ID if the image is any of the K persons (or “not recognized”)

1:K

K=100 ←



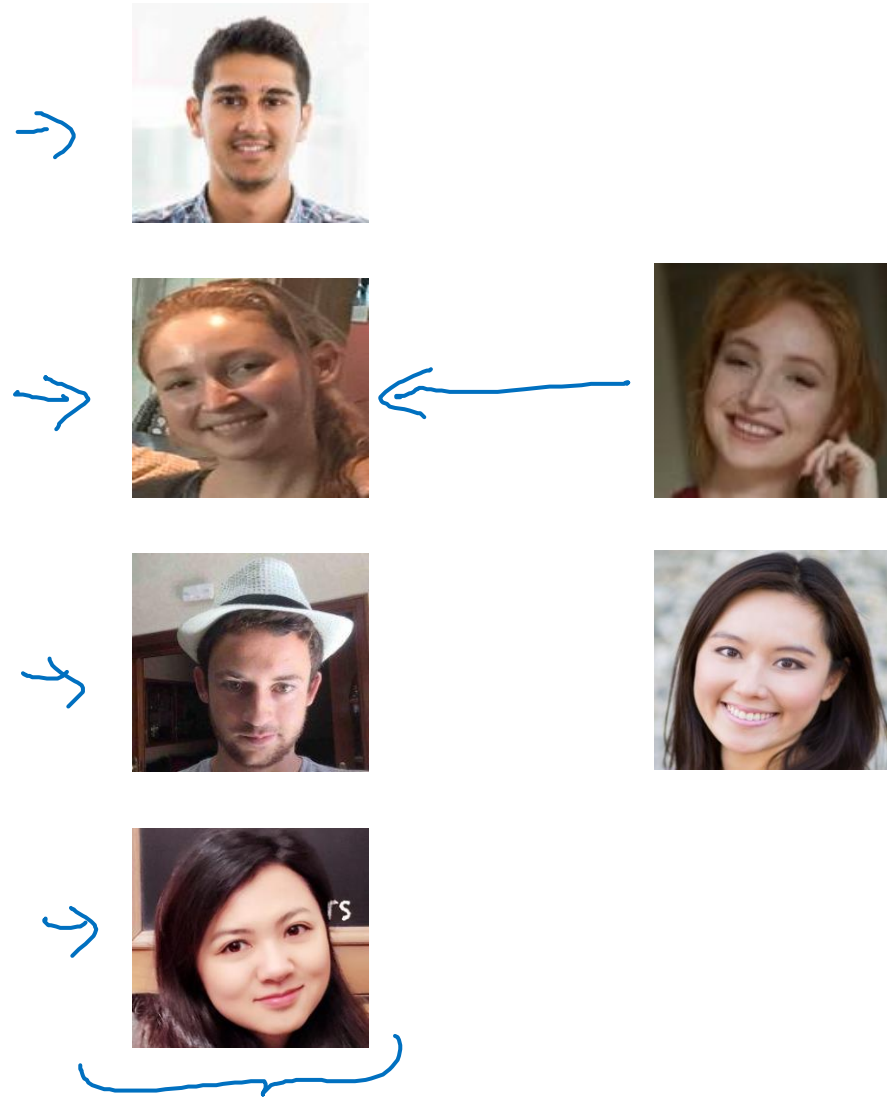
deeplearning.ai

# Face recognition

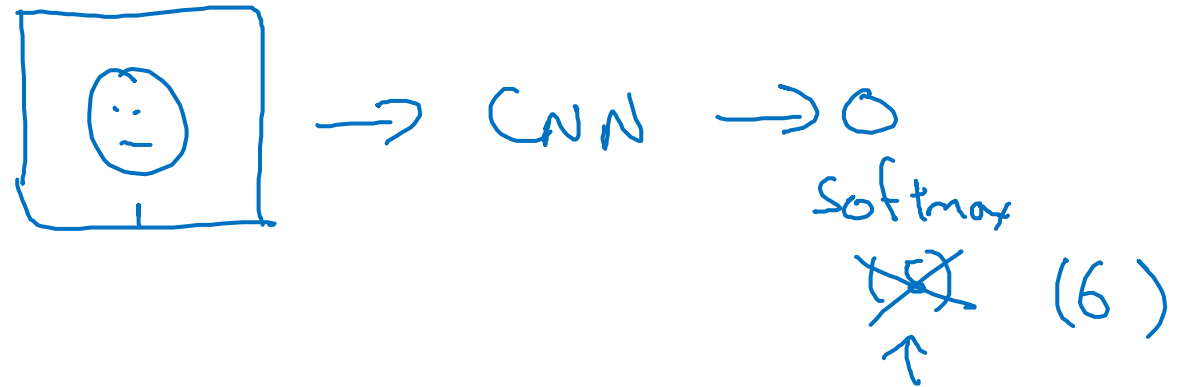
---

# One-shot learning

# One-shot learning



Learning from one example to recognize the person again



# Learning a “similarity” function

→  $d(\text{img1}, \text{img2})$  = degree of difference between images

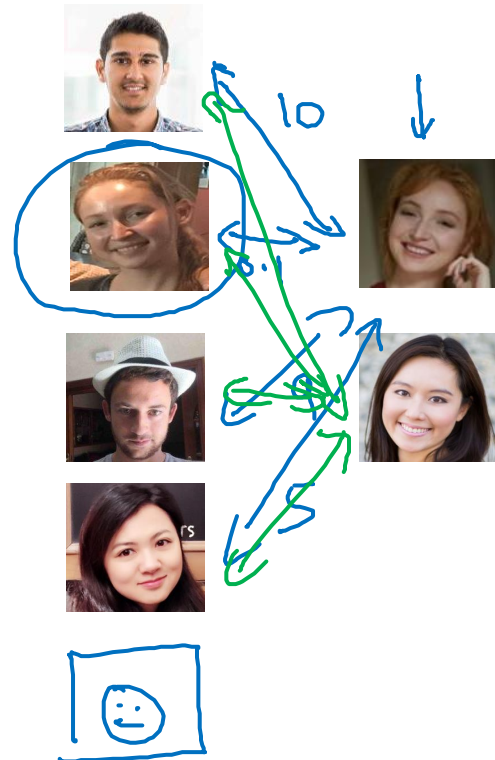
If  $d(\text{img1}, \text{img2}) \leq \tau$

$> \tau$

“same”

“different”

} Verification.



$d(\text{img1}, \text{img2})$



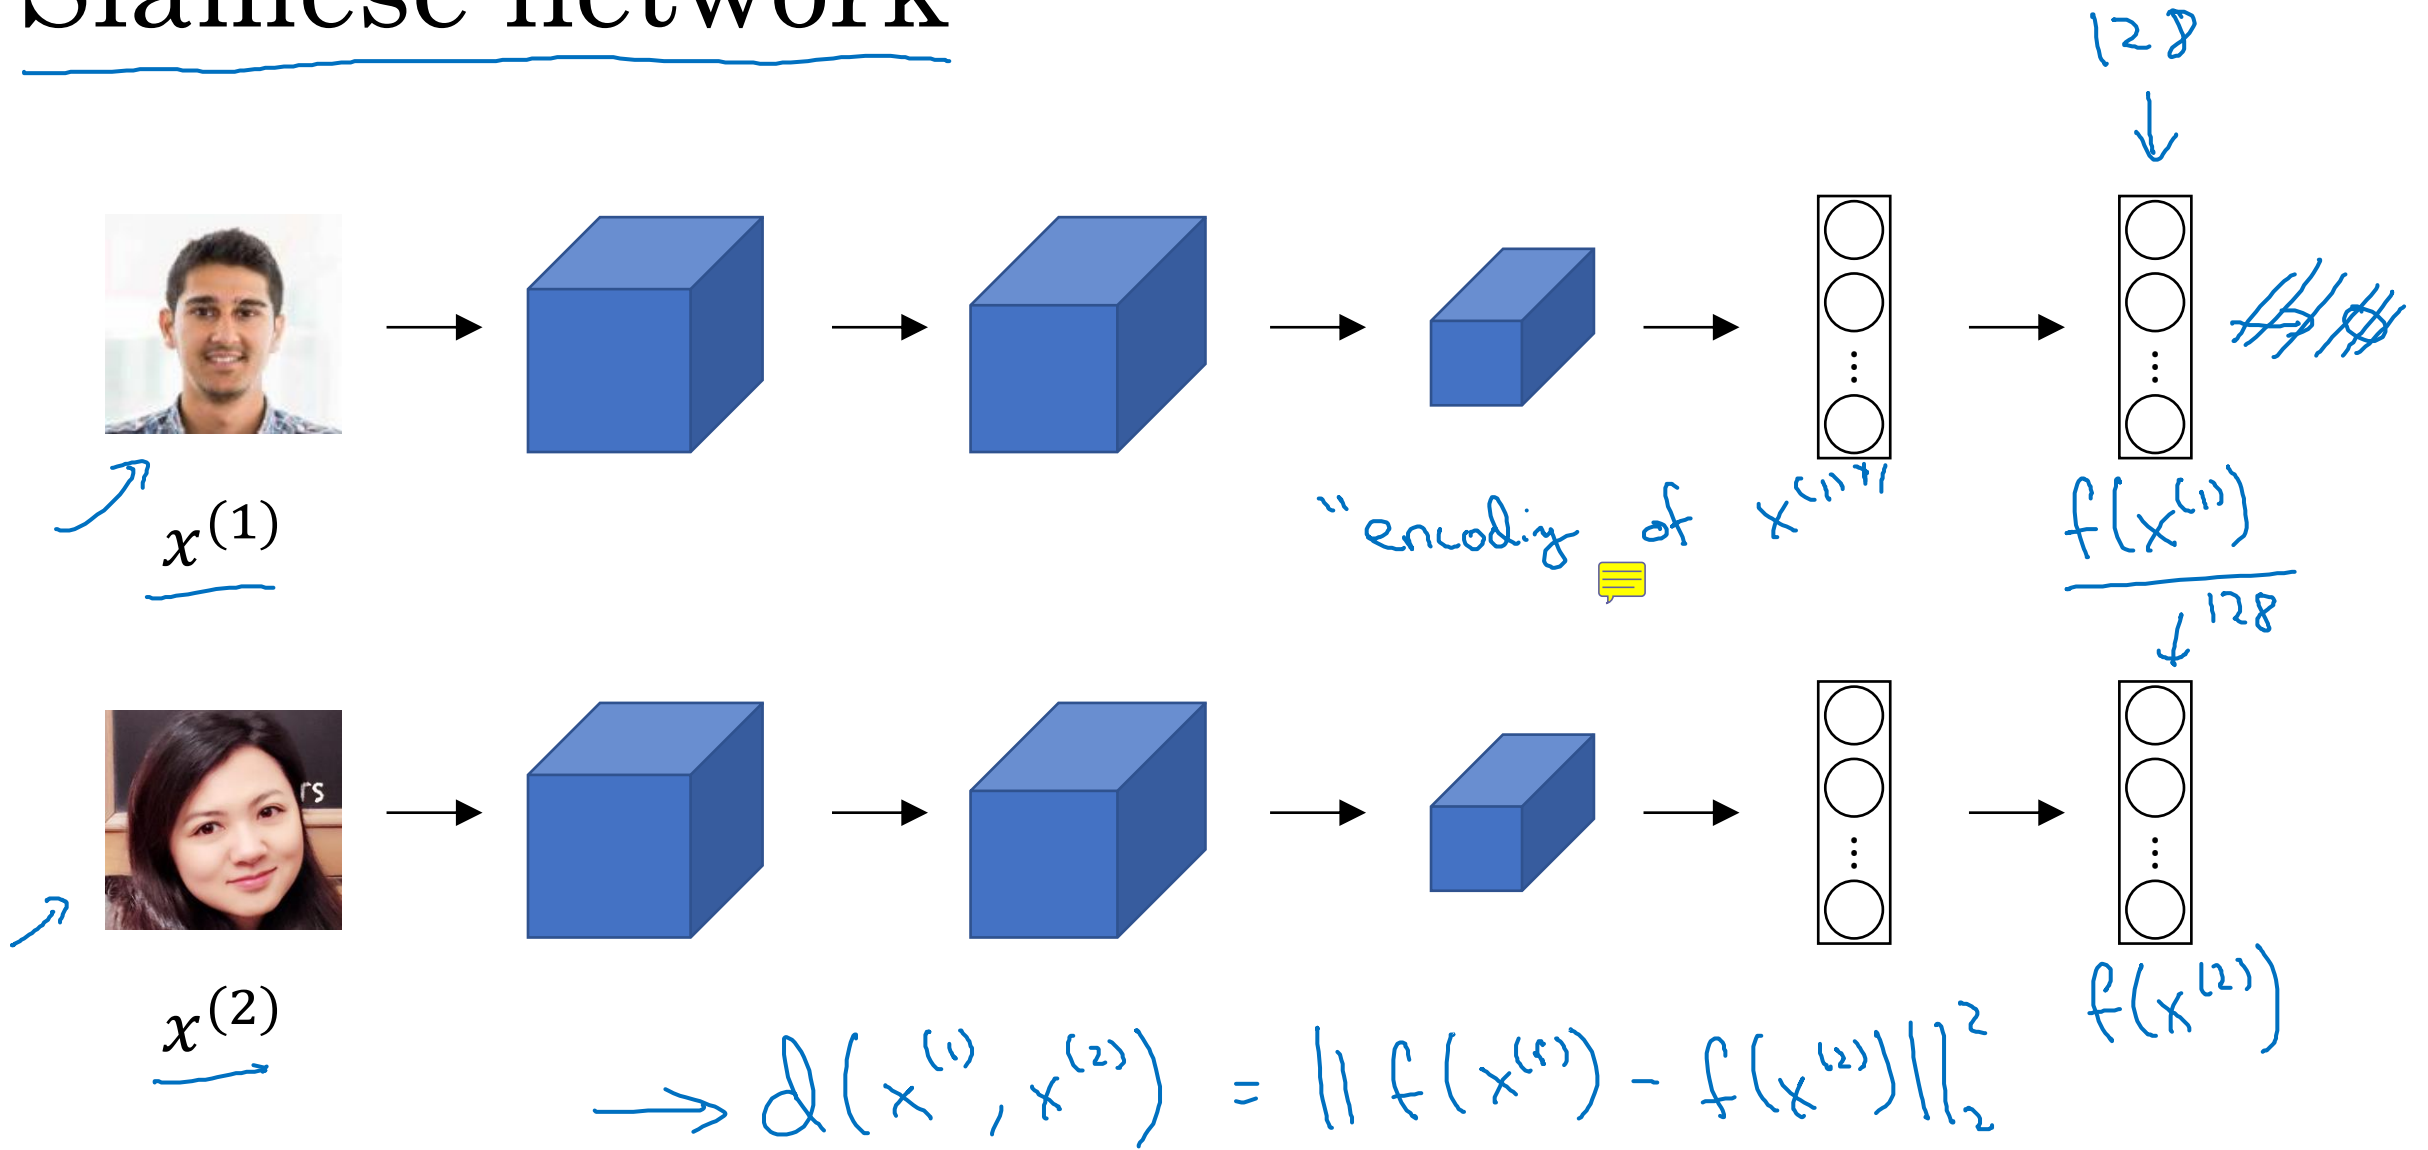
deeplearning.ai

# Face recognition

---


# Siamese network

# Siamese network





# Goal of learning



Parameters of NN define an encoding  $f(x^{(i)})$

Learn parameters so that:

If  $x^{(i)}, x^{(j)}$  are the same person,  $\|f(x^{(i)}) - f(x^{(j)})\|^2$  is small.

If  $x^{(i)}, x^{(j)}$  are different persons,  $\|f(x^{(i)}) - f(x^{(j)})\|^2$  is large.



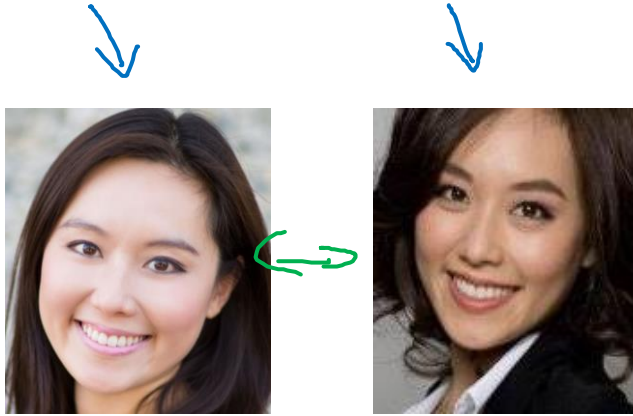
deeplearning.ai

# Face recognition

---

## Triplet loss

# Learning Objective



Anchor

A


Positive

P

$$d(A, P) = 0.5$$

Want:

$$\underbrace{\|f(A) - f(P)\|^2}_{d(A, P)} + \underline{\alpha} \leq \overset{\nearrow 0.2}{\phantom{\alpha}}$$



Anchor

A

Negative

N

$$d(A, N) = \cancel{0.5} \quad 0.7$$

$$\underbrace{\|f(A) - f(N)\|^2}_{d(A, N)}$$

$$\underbrace{\|f(A) - f(P)\|^2}_0 - \underbrace{\|f(A) - f(N)\|^2}_0 + \underline{\alpha} \leq \underline{0} \quad \text{margin}$$

$$f(\text{img}) = \vec{0}$$

# Loss function

Given 3 images

$A, P, N$ :

$$\underline{L(A, P, N)} = \max \left( \underbrace{\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha}_{\geq 0}, 0 \right)$$

$$J = \sum_{i=1}^m L(A^{(i)}, P^{(i)}, N^{(i)})$$

$A, P$   
 $\uparrow \quad \uparrow$

Training set: 10k pictures of 1k persons

# Choosing the triplets A,P,N

During training, if A,P,N are chosen randomly,  
 $d(A, P) + \alpha \leq d(A, N)$  is easily satisfied.

$$\|f(A) - f(P)\|^2 + \alpha \leq \|f(A) - f(N)\|^2$$

Choose triplets that're "hard" to train on.

$$\frac{d(A, P)}{d(A, P)} + \alpha \leq \frac{d(A, N)}{d(A, N)}$$

↓                      ↑

Face Net  
Deep Face

# Training set using triplet loss

Anchor



⋮



Positive




⋮



Negative



⋮



J

$$d(x^{(i)}, x^{(j)})$$




deeplearning.ai

# Face recognition

---









## Face verification and binary classification

# Learning the similarity function





# Face verification supervised learning

$x$		$y$	
		1	"Same"
		0	"Different"
		0	
		1	



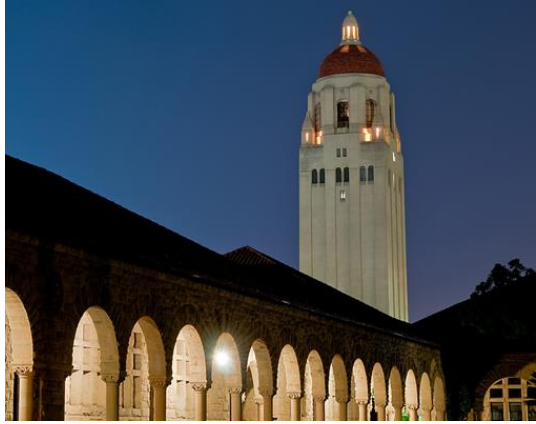
deeplearning.ai

# Neural Style Transfer

---

What is neural style  
transfer?


# Neural style transfer



Content ( $c$ )



Style ( $s$ )



Generated image ( $G$ )



Content ( $c$ )



Style ( $s$ )



Generated image ( $G$ )




deeplearning.ai

# Neural Style Transfer

---

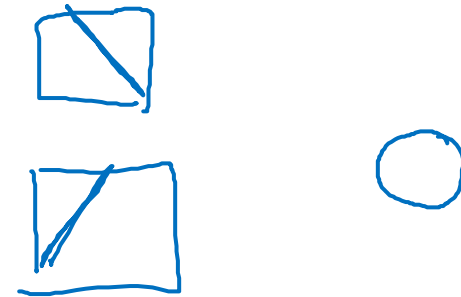
What are deep  
ConvNets learning?

# Visualizing what a deep network is learning



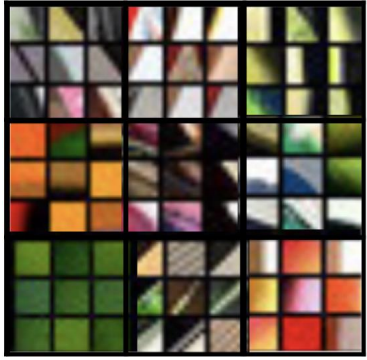
Pick a unit in layer 1. Find the nine image patches that maximize the unit's activation.

Repeat for other units.






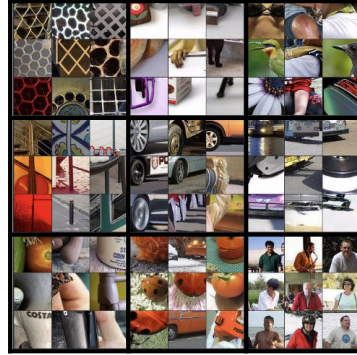
# Visualizing deep layers



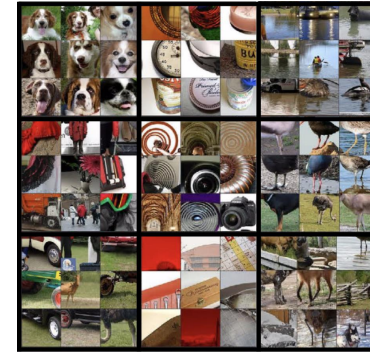
Layer 1



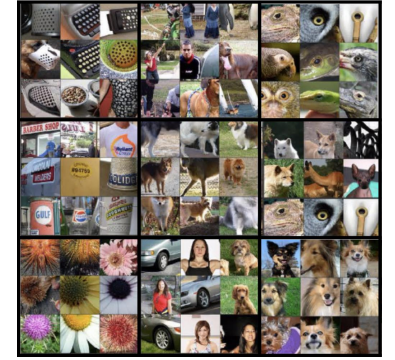
Layer 2



Layer 3

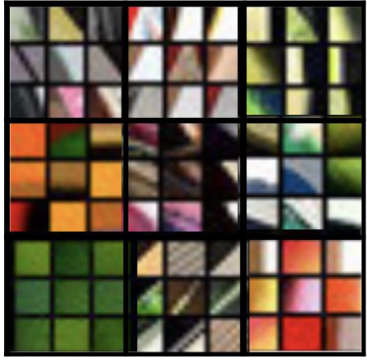


Layer 4

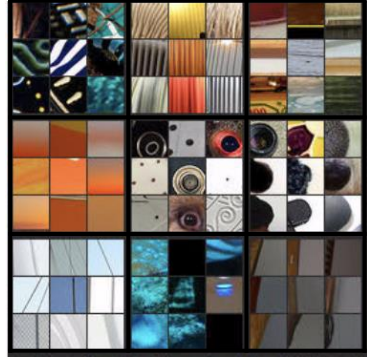


Layer 5

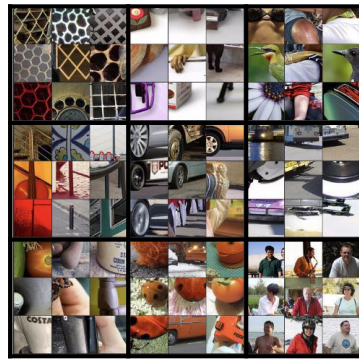
# Visualizing deep layers: Layer 1



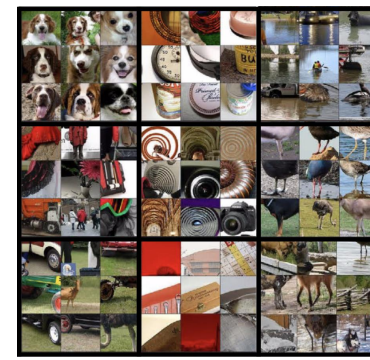
Layer 1



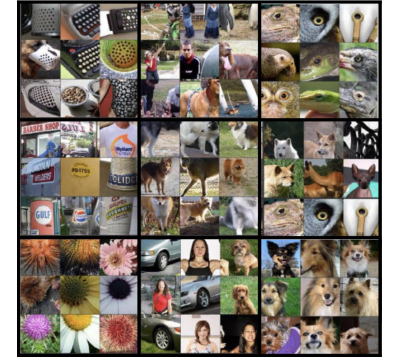
Layer 2




Layer 3



Layer 4

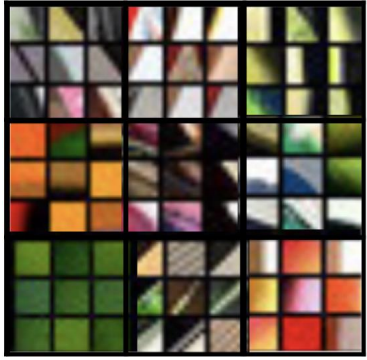


Layer 5

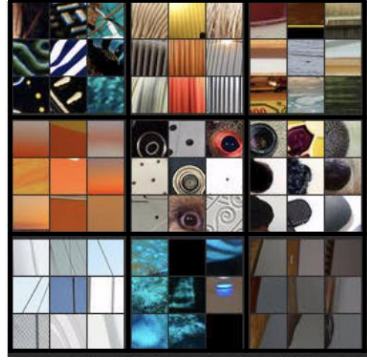




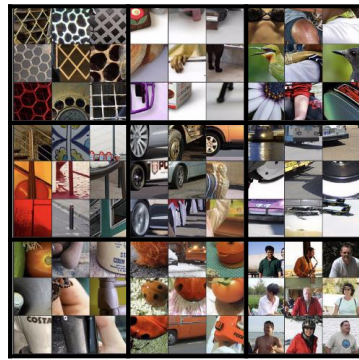
# Visualizing deep layers: Layer 2



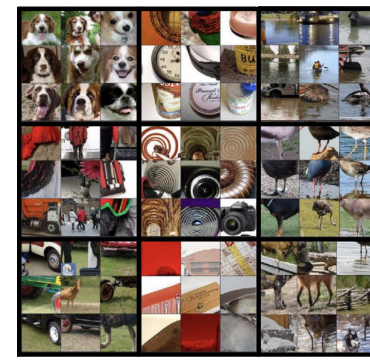
Layer 1



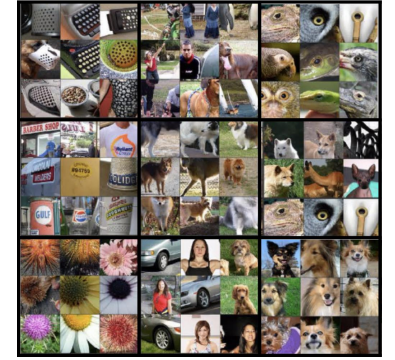
Layer 2



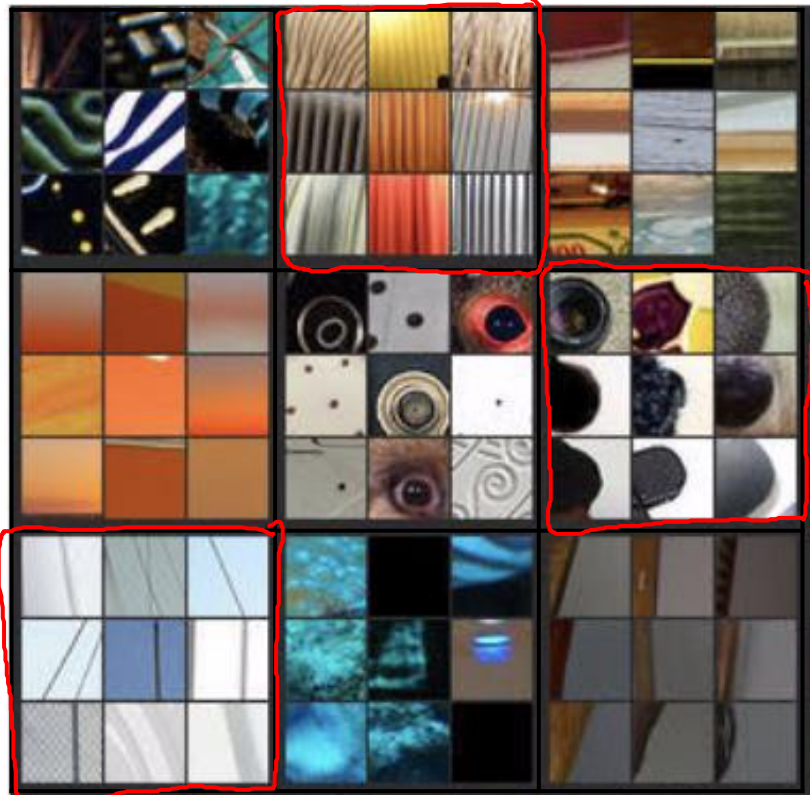
Layer 3



Layer 4

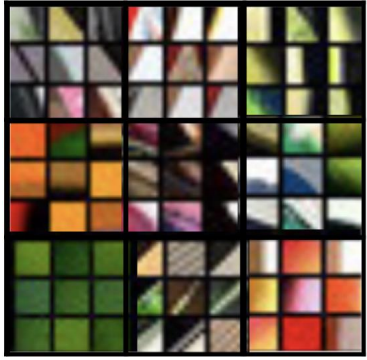


Layer 5







# Visualizing deep layers: Layer 3




Layer 1




Layer 2




Layer 3



Layer 4

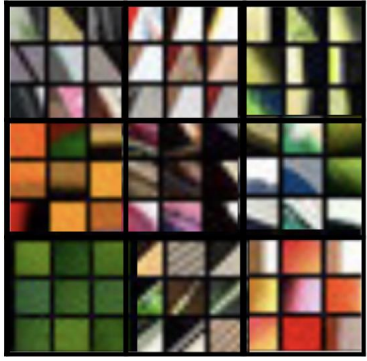


Layer 5






# Visualizing deep layers: Layer 3

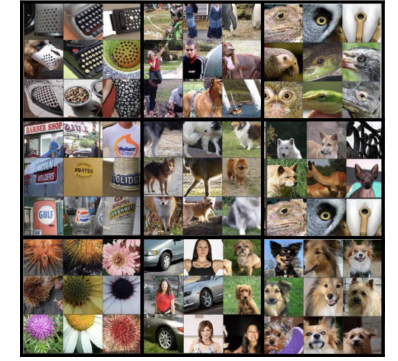


Layer 1



Layer 3


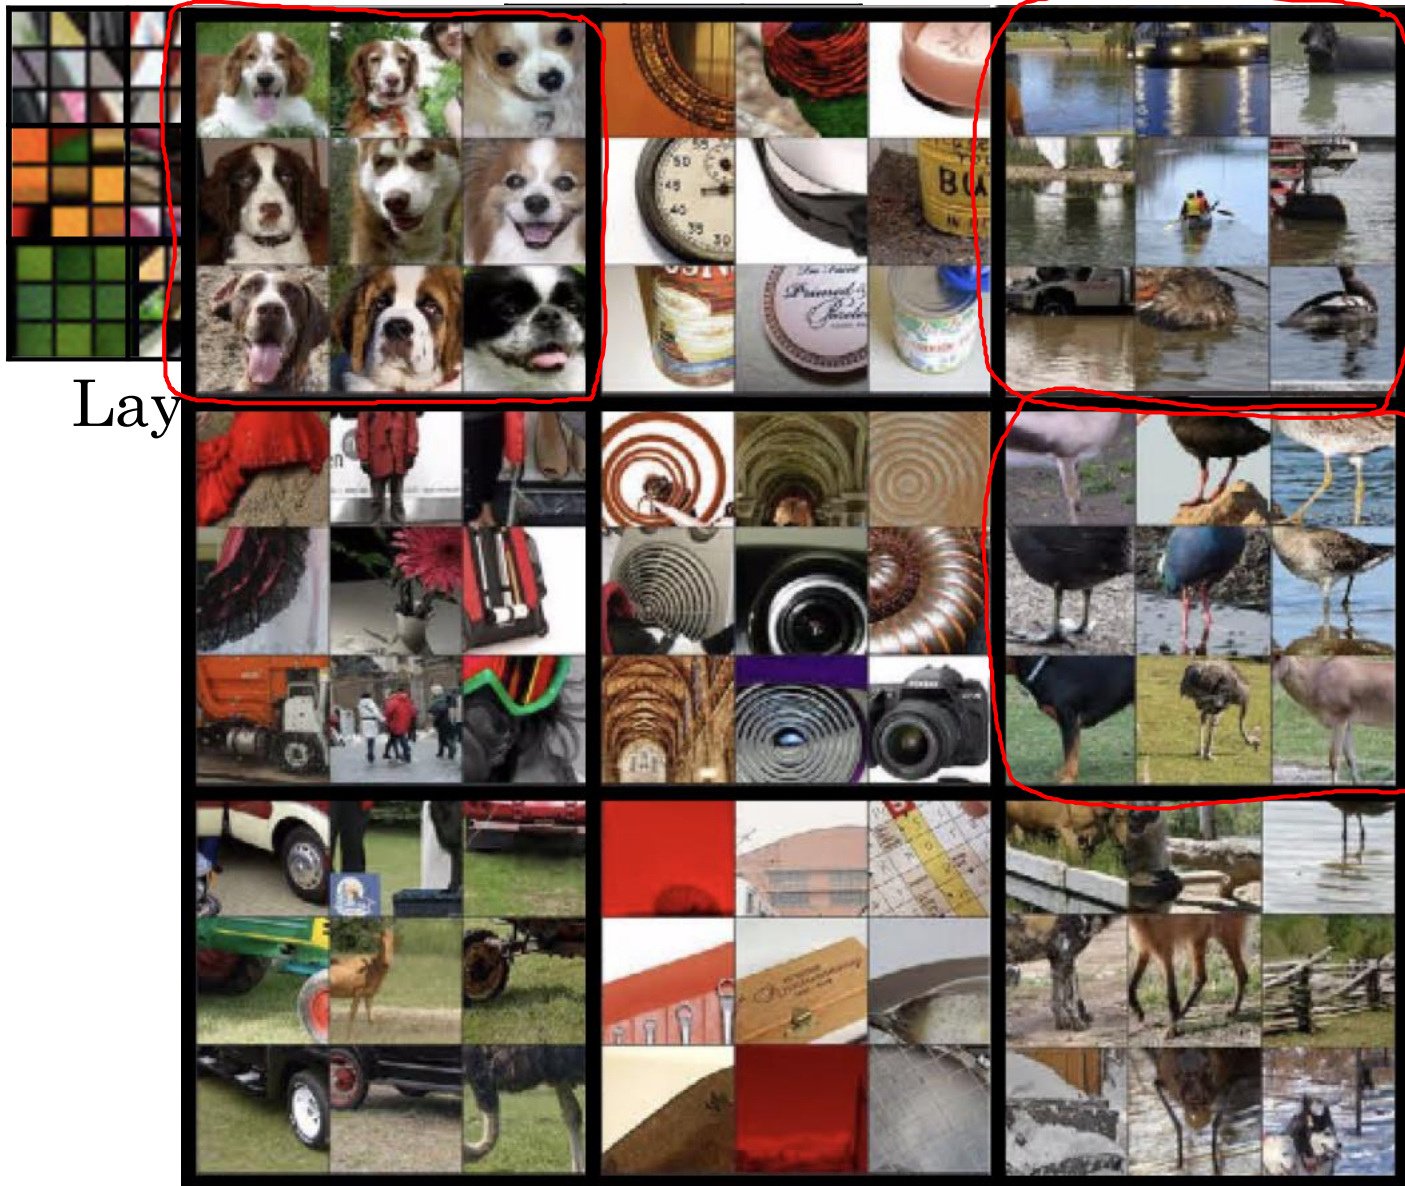
Layer 4




Layer 5



# Visualizing deep layers: Layer 4



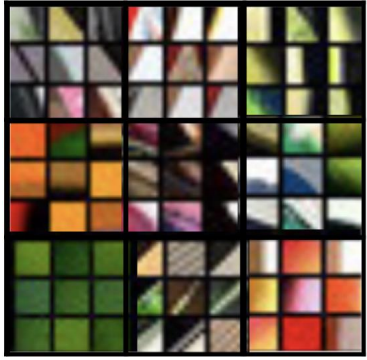
Layer 4



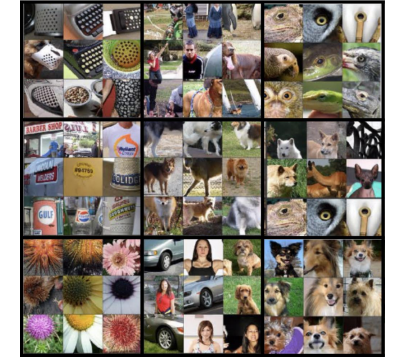

Layer 5



# Visualizing deep layers: Layer 5



Layer 1



Layer 5



deeplearning.ai

# Neural Style Transfer

---

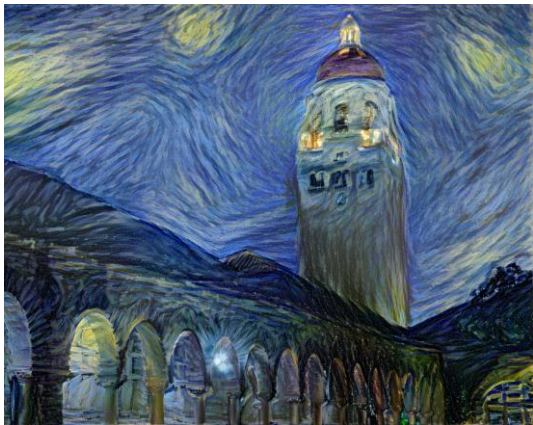
## Cost function

# Neural style transfer cost function



Content C

Style S



Generated image G

$$\mathcal{J}(G) = \alpha \mathcal{J}_{\text{Content}}(C, G) + \beta \mathcal{J}_{\text{Style}}(S, G)$$



Find the generated image  $G$

- ## 1. Initiate $G$ randomly

G:  $\underline{100} \times \underline{100} \times \underset{\substack{\uparrow \\ RGB}}{3}$

2. Use gradient descent to minimize  $J(G)$

$$G := G - \frac{d}{2G} J(G)$$





deeplearning.ai

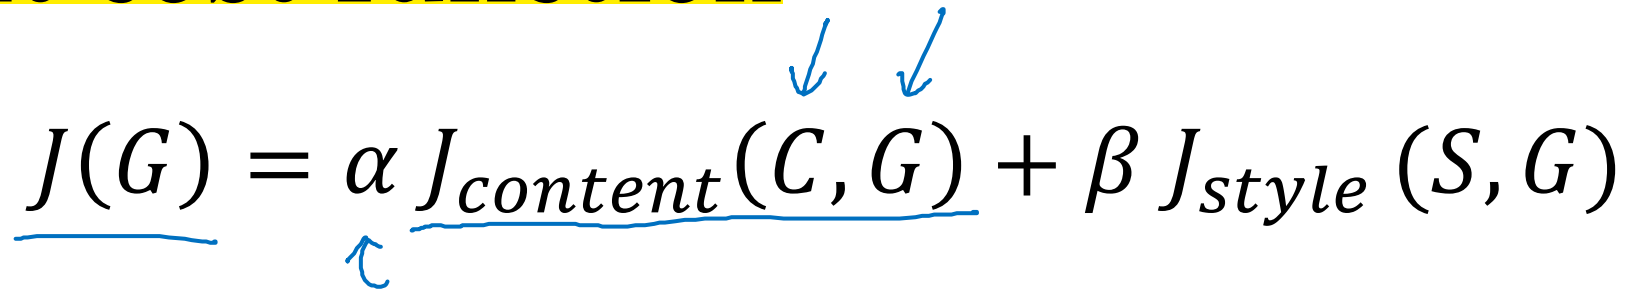
# Neural Style Transfer

---

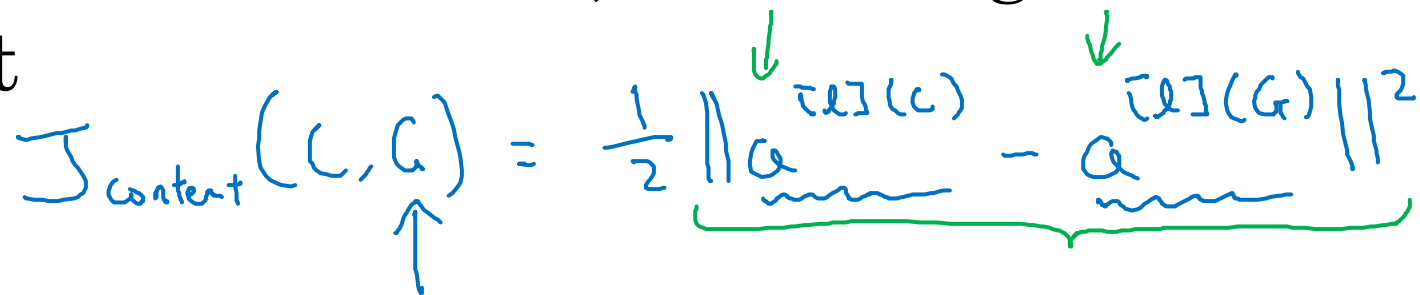
## Content cost function



# Content cost function

$$\underline{J(G)} = \alpha \underline{J_{content}(C, G)} + \beta J_{style}(S, G)$$


- Say you use hidden layer  $l$  to compute content cost.
- Use pre-trained ConvNet. (E.g., VGG network)
- Let  $a^{[l](C)}$  and  $a^{[l](G)}$  be the activation of layer  $l$  on the images
- If  $a^{[l](C)}$  and  $a^{[l](G)}$  are similar, both images have similar content

$$J_{content}(C, G) = \frac{1}{2} \left\| \underbrace{a^{[l](C)}}_{\text{green arrow}} - \underbrace{a^{[l](G)}}_{\text{green arrow}} \right\|^2$$





deeplearning.ai

# Neural Style Transfer


---

## Style cost function

# Meaning of the “style” of an image




Say you are using layer  $l$ 's activation to measure “style.”  
Define style as correlation between activations across channels.




How correlated are the activations  
across different channels?

# Intuition about style of an image


Style image



Generated Image



Correlated?  
Uncorrelated



# Style matrix

Let  $a_{i,j,k}^{[l]}$  = activation at  $(i, j, k)$ .  $\underline{G}^{[l]}$  is  $\underline{n}_c^{[l]} \times \underline{n}_c^{[l]}$

H W C  
↓ ↓ ↙

$n_c$   
 $G_{kk'}^{[l]}$   
↑ ↑  
 $k = 1, \dots, n_c$

style of Style Image

→  $G_{kk'}^{[l]}(S)$

→  $G_{kk'}^{[l]}(G)$

style of Generated Image

$$G_{kk'}^{[l]}(S) = \sum_{i=1}^{n_H^{[l]}} \sum_{j=1}^{n_W^{[l]}} a_{ijk}^{[l]}(S) a_{ijk'}^{[l]}(S)$$

$$G_{kk'}^{[l]}(G) = \sum_{i=1}^{n_H^{[l]}} \sum_{j=1}^{n_W^{[l]}} a_{ijk}^{[l]}(G) a_{ijk'}^{[l]}(G)$$

"Gram matrix"

$$J_{\text{style}}^{[l]}(S, G) = \frac{1}{(\dots)} \left\| G^{[l]}(S) - G^{[l]}(G) \right\|_F^2$$

$$\stackrel{\beta}{\uparrow} = \frac{1}{(2 n_H^{[l]} n_W^{[l]} n_c^{[l]})^2} \sum_k \sum_{k'} (G_{kk'}^{[l]}(S) - G_{kk'}^{[l]}(G))^2$$

# Style cost function

$$\|G^{[l](S)} - G^{[l](G)}\|_F^2$$

$$J_{style}^{[l]}(S, G) = \frac{1}{\left(2n_H^{[l]}n_W^{[l]}n_C^{[l]}\right)^2} \sum_k \sum_{k'} (G_{kk'}^{[l](S)} - G_{kk'}^{[l](G)})^2$$

$$J_{style}(S, G) = \sum_l \lambda_l J_{style}^{[l]}(S, G)$$

$$\underbrace{J(G)}_G = \alpha J_{content}(G) + \beta J_{style}(S, G)$$




deeplearning.ai

# Convolutional Networks in 1D or 3D

---

1D and 3D  
generalizations of  
models

# Convolutions in 2D and 1D



$$14 \times 14 \times \underline{3} * 5 \times 5 \times \underline{3}$$

$$\rightarrow \underline{10 \times 10 \times 16}$$

$$\underline{10 \times 10 \times 16} * \underline{5 \times 5 \times 16}$$

$$\rightarrow \underline{6 \times 6 \times 32}$$

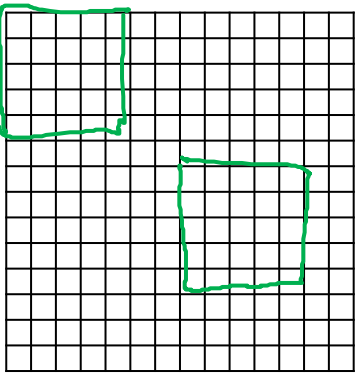
---

$$14 \times \underline{1} * 5 \times \underline{1}$$

$$\rightarrow 10 \times \underline{16}$$

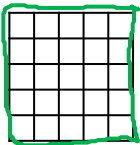
$$\underline{10 \times 16} * \underline{5 \times 16}$$

$$\rightarrow \underline{6 \times 32}$$




2D input image  
14 × 14



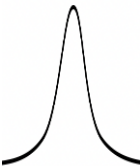
\*



2D filter  
5 × 5




\*

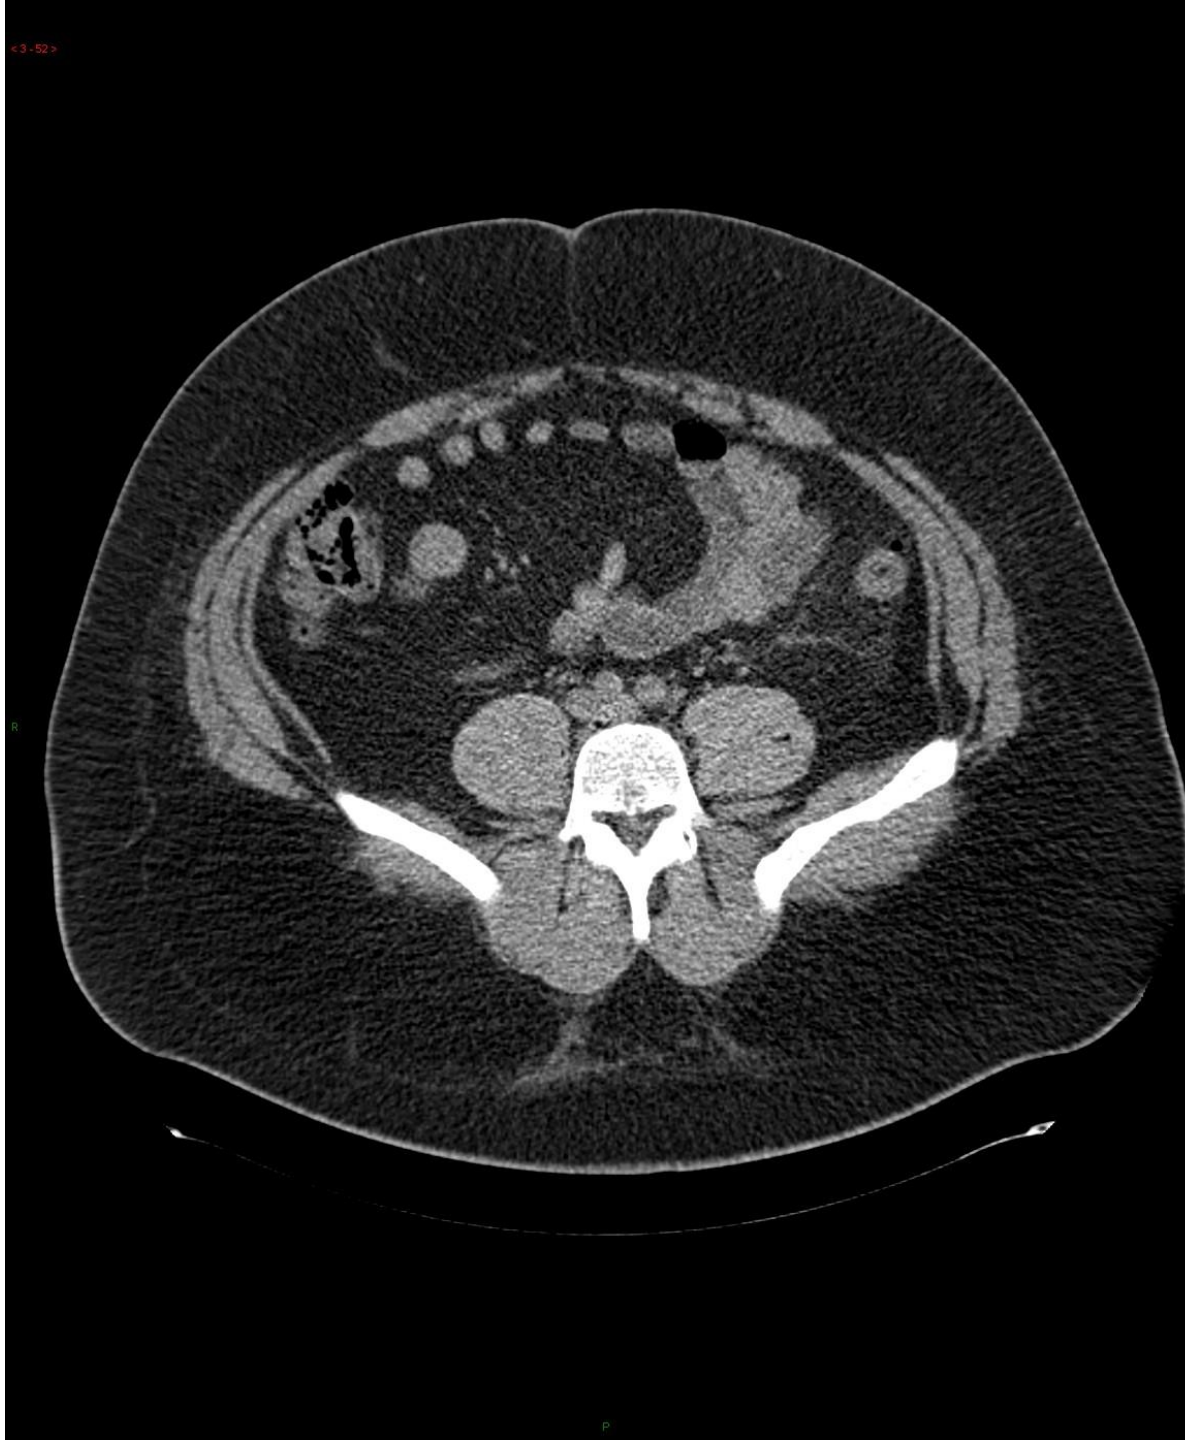




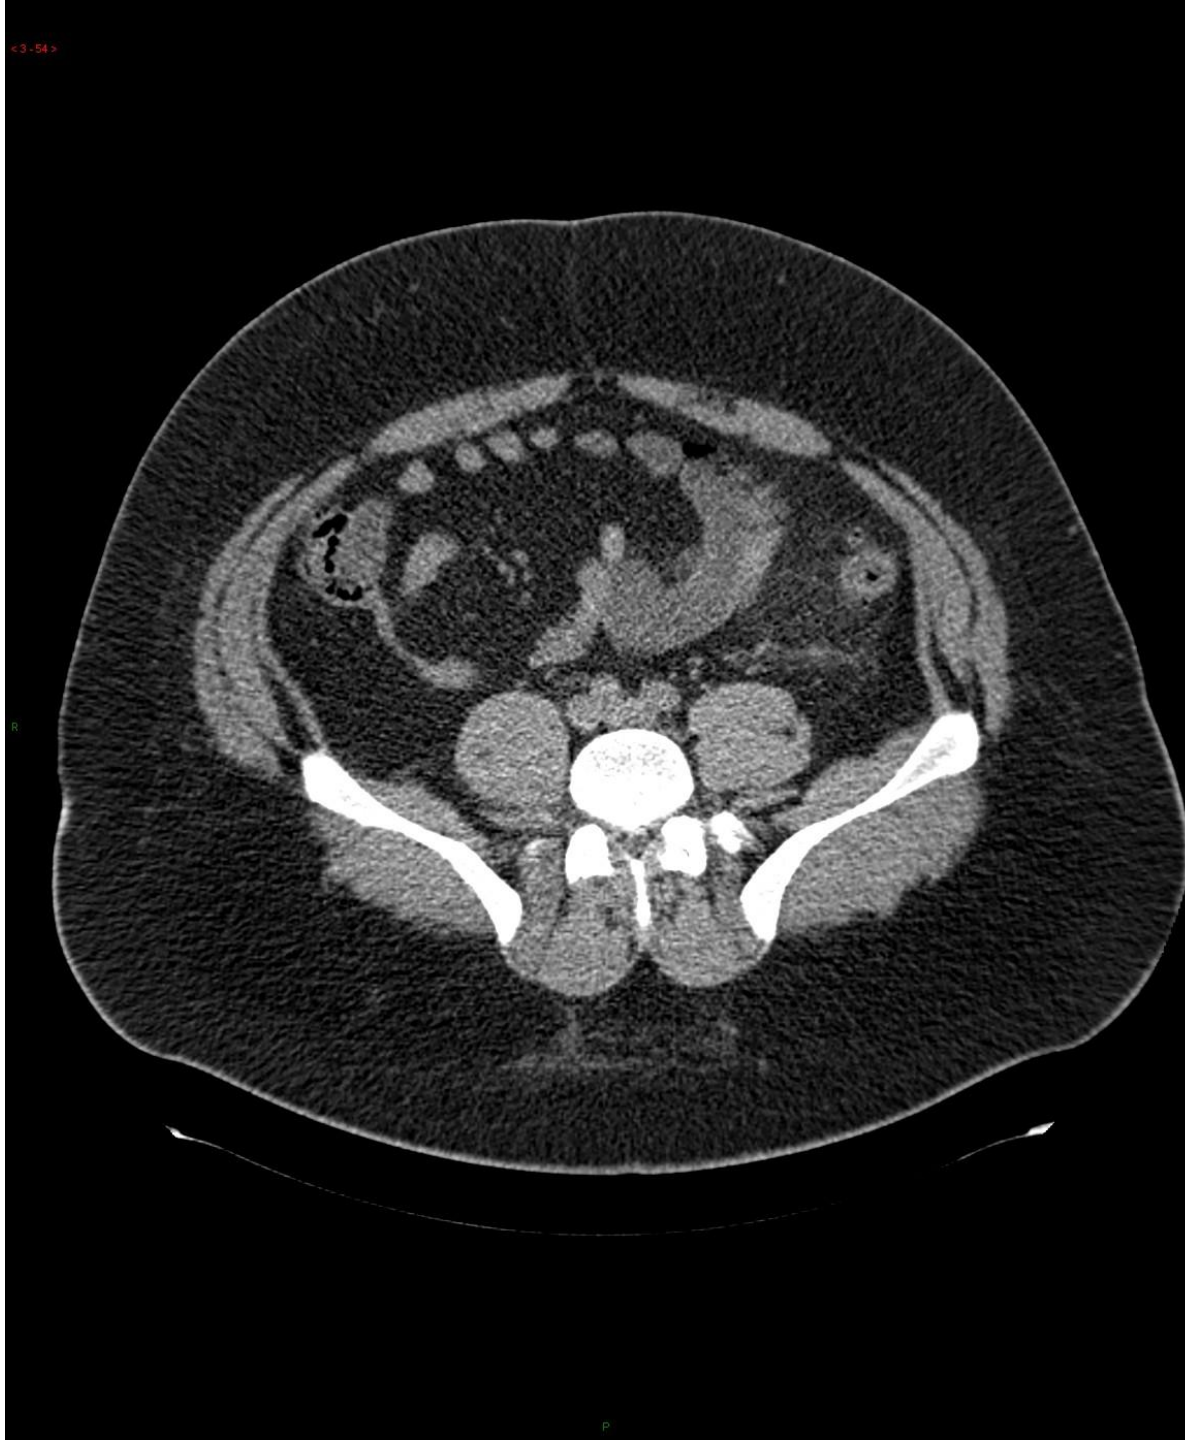
3D data



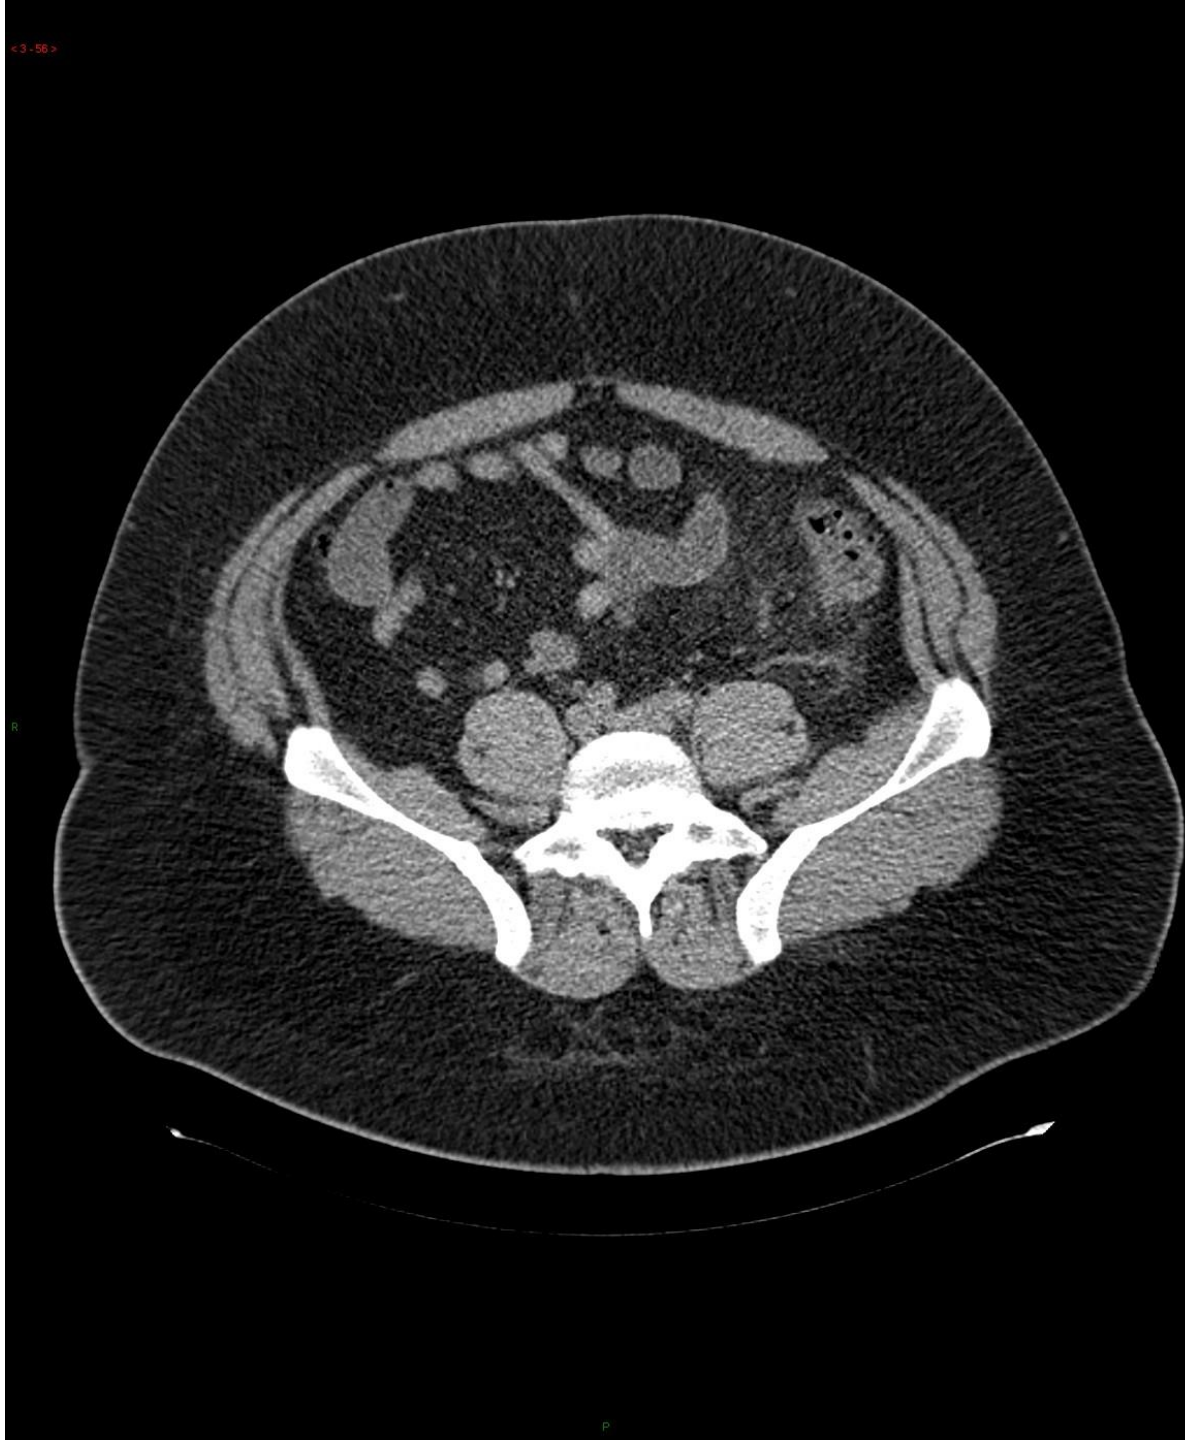
3D data



3D data

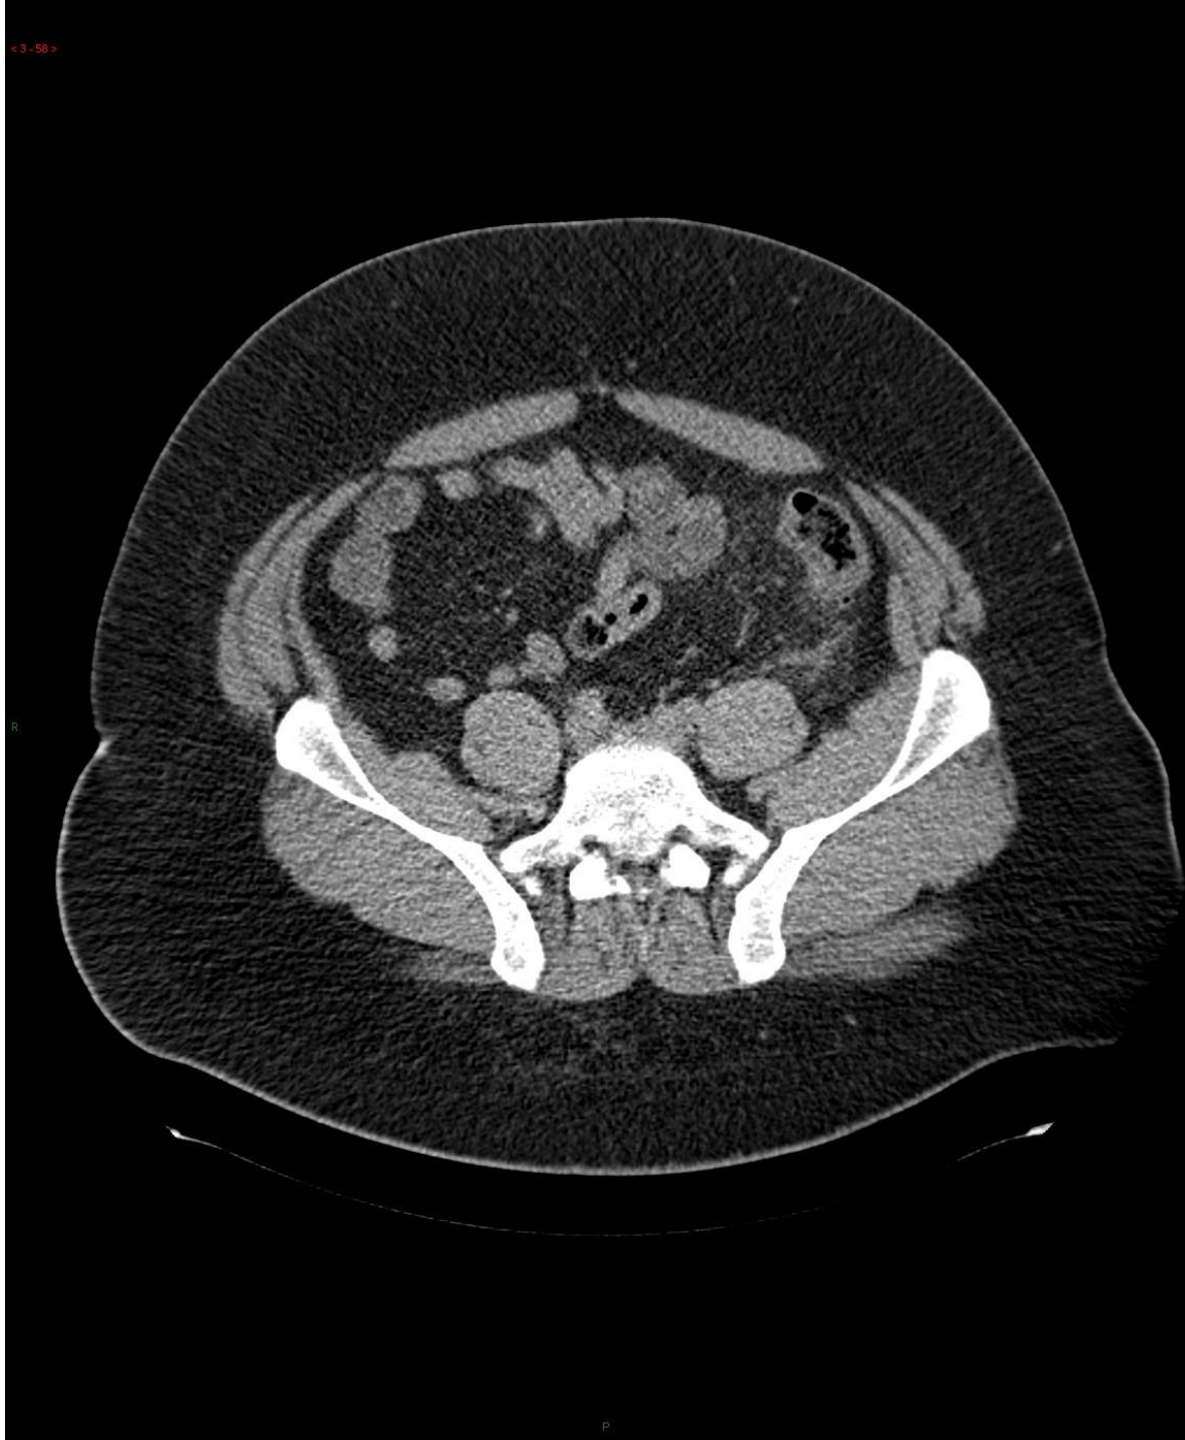


3D data





# 3D data



# 3D data



# 3D data



# 3D data





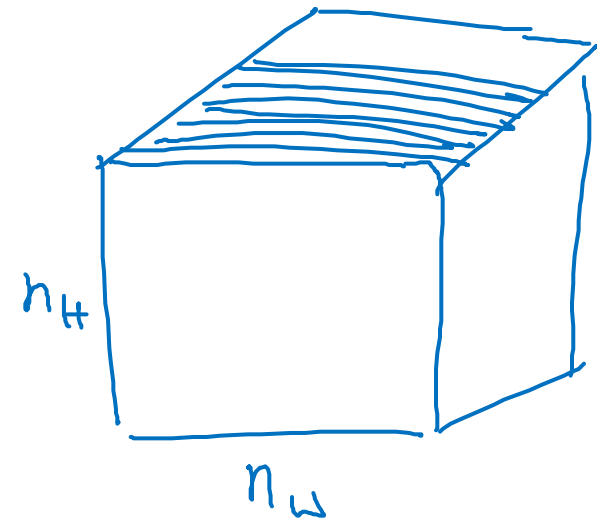
# 3D data




3D data




# 3D data



# 3D convolution



3D volume



$$\begin{array}{l}
 \downarrow \quad \downarrow \quad \downarrow \quad \downarrow^{n_c} \\
 \underline{14 \times 14 \times 14} \times \underline{1} \\
 * \quad \underline{5 \times 5 \times 5} \times \underline{1} \quad 16 \text{ filters} \\
 \rightarrow 10 \times 10 \times 10 \times \underline{16} \\
 * 5 \times 5 \times 5 \times \underline{16} \quad 32 \text{ filters} \\
 \rightarrow 6 \times 6 \times 6 \times 32
 \end{array}$$